


Visual Intelligence Project

Lung Cancer Histopathological Classification

Your Name | March 2025

🔍 Introduction & Problem Statement

- **Dataset:** Lung cancer histopathological images (3 classes):
 - Adenocarcinoma
 - Squamous cell carcinoma
 - Benign tissue
- **Classification Task:** Binary classification (adenocarcinoma vs benign)
- **Challenge:** Distinguishing subtle tissue patterns and cellular structures





Project Goals

- Compare traditional CNN vs ScatNet approaches
- Investigate color vs structural features
- Achieve high accuracy with interpretable results
- Apply explainable AI techniques to validate model decisions



Data Preprocessing & Setup

- **Dataset Organization:**
 - K-fold cross-validation with 10 folds
 - Balanced class distribution
 - Target size: 768×768 pixels (original)
→ 224×224 (processed)
- **Key Finding:** The class average color alone is sufficient for high-accuracy classification
 - This prompted our focus on structural features identification





Preprocessing Pipeline


- Image normalization and standardization
- Data augmentation decisions:
 - Rotations ($0^\circ, 90^\circ, 180^\circ, 270^\circ$)
 - Horizontal and vertical flips
 - Minor elastic distortions
- Color vs grayscale analysis
 - With grayscale images, models can still achieve high accuracy
 - Focus shifted to structural features identification

Key Discovery: Models heavily rely on color features for classification, but structural features are critical for generalization



CNN Model Architecture

- Efficient convolutional neural network:
 - Two convolutional blocks with batch normalization
 - Input channels: 1 (grayscale) or 3 (RGB)
 - Simple classifier head (16-dimensional feature space)
 - **Key Finding:** Good at learning even with a small classifier layer



```
graph LR  
    A[Input Image] --> B[Conv Blocks] --> C[Global Pooling] --> D[Classifier] --> E[Output]
```




CNN Architecture Details

Layer (type:depth-idx)	Output Shape	Param #
CNNImageClassifier	[1, 2]	--
Sequential: 1-1	[1, 24, 4, 4]	--
Conv2d: 2-1	[1, 16, 382, 382]	1,952
BatchNorm2d: 2-2	[1, 16, 382, 382]	32
ReLU: 2-3	[1, 16, 382, 382]	--
MaxPool2d: 2-4	[1, 16, 191, 191]	--
Conv2d: 2-5	[1, 16, 191, 191]	2,320
BatchNorm2d: 2-6	[1, 16, 191, 191]	32
ReLU: 2-7	[1, 16, 191, 191]	--
MaxPool2d: 2-8	[1, 16, 95, 95]	--
Conv2d: 2-9	[1, 24, 95, 95]	3,480
BatchNorm2d: 2-10	[1, 24, 95, 95]	48
ReLU: 2-11	[1, 24, 95, 95]	--
AdaptiveAvgPool2d: 2-12	[1, 24, 4, 4]	--
FeatureClassifier: 1-2	[1, 2]	--
Linear: 2-13	[1, 16]	6,160
BatchNorm1d: 2-14	[1, 16]	32
ReLU: 2-15	[1, 16]	--
Dropout: 2-16	[1, 16]	--
Linear: 2-17	[1, 2]	34



ScatNet Model Architecture

- Wavelet-based feature extraction:
 - $J=3$ scale parameter for wavelet decomposition
 - $L=8$ orientations, $M=2$ scattering order
 - Complex classifier ($217 \rightarrow 64 \rightarrow 2$ neurons)
 - Translation, rotation, and scaling invariant
- **Key Finding:** Requires more complex classifier layer to achieve good performance





Training & Evaluation Metrics

Metric	CNN	ScatNet
Mean Accuracy	98.9%	87.9%
Mean F1 Score	98.9%	86.7%
Best Fold Accuracy	99.6%	92.6%
Training Speed	Faster	Slower
Classifier Complexity	Simple	Complex

Note: Values directly from the 10-fold cross-validation experiments



CNN Performance Details

Fold-by-fold performance:

- Fold 0: 97.6% Accuracy, 97.6% F1 Score
- Fold 1: 97.6% Accuracy, 97.6% F1 Score
- Fold 2: 99.2% Accuracy, 99.2% F1 Score
- Fold 3: 99.3% Accuracy, 99.3% F1 Score
- Fold 4: 98.2% Accuracy, 98.2% F1 Score
- Fold 5: 99.3% Accuracy, 99.3% F1 Score
- Fold 6: 99.5% Accuracy, 99.5% F1 Score
- Fold 7: 99.1% Accuracy, 99.1% F1 Score
- Fold 8: 99.6% Accuracy, 99.6% F1 Score
- Fold 9: 99.3% Accuracy, 99.3% F1 Score



ScatNet Performance Details

Fold-by-fold performance:

- Fold 0: 91.1% Accuracy, 90.8% F1 Score
- Fold 1: 92.6% Accuracy, 92.7% F1 Score
- Fold 2: 88.7% Accuracy, 87.8% F1 Score
- Fold 3: 88.9% Accuracy, 88.0% F1 Score
- Fold 4: 88.2% Accuracy, 87.3% F1 Score
- Fold 5: 81.1% Accuracy, 77.4% F1 Score (lowest)
- Fold 6: 85.0% Accuracy, 82.8% F1 Score
- Fold 7: 87.8% Accuracy, 86.7% F1 Score
- Fold 8: 88.9% Accuracy, 88.2% F1 Score
- Fold 9: 86.5% Accuracy, 85.2% F1 Score





Key Performance Findings

- CNN significantly outperforms ScatNet in both accuracy and speed
 - 11% mean accuracy difference (98.9% vs 87.9%)
- K-fold validation confirms robust performance across data splits
 - CNN shows less variance between folds
- CNN achieves convergence in fewer epochs
- Performance gap indicates CNN's superior ability to learn relevant features

CNN Learning Curves Analysis



- Rapid convergence within 10-15 epochs
- Consistent performance across folds
- Limited overfitting due to effective regularization
- Final validation accuracy stabilized around 99%
- More efficient training compared to ScatNet



ScatNet Learning

Curves Analysis


- Slower convergence requiring more epochs
- Higher variance between folds (81.1% - 92.6%)
- More complex classifier needed for good performance
- Validation accuracy plateaued around 88%
- Greater performance variation across data splits





CNN Filter Analysis


- Learned color-sensitive patterns automatically
- Hierarchical feature extraction with progressive abstraction
- First layer captures basic edges and textures
- Deeper layers identify tissue-specific patterns
- Impact of data augmentation: improved filter robustness to orientations



🔬 ScatNet Filter Analysis


- Pre-defined wavelet transforms (not learned)
- Scale and rotation invariant features
- Lower discriminative power despite theoretical advantages
- Fixed mathematical representation limits adaptability
- Data augmentation impact: less significant due to inherent invariance

Wavelets for each scales j and angles θ used.
Color saturation and color hue respectively denote complex magnitude and complex phase.



Explainable AI Methods


- **Custom Attribution Methods:**
 - Vanilla Backpropagation
 - Guided Backpropagation
 - Occlusion
- **Captum Library Integration:**
 - Multiple attribution techniques
 - Heatmap visualization highlighting decision regions





CNN Attribution Analysis


- Visualizes regions most influential for classification decisions
- Focuses on cellular structures and color patterns
- Higher resolution in feature attribution
- Strong correlation between attribution maps and pathological markers





ScatNet Attribution Analysis

- Different activation patterns compared to CNN
- More diffuse attribution regions
- Wavelets capture texture but miss important color information
- Less aligned with pathological indicators



Custom vs. Library Implementation

- **Custom Implementation:**

- Complete control over visualization parameters
- Direct access to gradient computation
- Greater understanding of attribution mechanics

- **Captum Library:**

- More visualization options and integrated smoothing
- Consistent API across different attribution methods
- Better computational performance

Analysis: Both implementations highlight similar regions, validating our approach



Performance Summary

- **Performance Achievements:**
 - CNN: 98.9% mean accuracy with simpler architecture
 - ScatNet: 87.9% mean accuracy despite theoretical advantages
 - 11% performance gap between approaches
- **Computational Efficiency:**
 - CNN showed faster training times
 - ScatNet required more complex classifier to achieve reasonable performance



Key Insights

- Color features are crucial for lung cancer histopathology classification
 - The class average color alone is highly predictive
- Learned features (CNN) outperform fixed mathematical representations (ScatNet)
- Simpler architectures can outperform sophisticated ones when aligned with data characteristics
- Even with grayscale images, the models achieved high accuracy
 - Focusing on structural features rather than just color

Thank You!

Questions?

 your.email@university.edu

 github.com/yourusername